

## Summary of changes:

Steps taken to disable KV caching:

1. In model.py modify Attention.forward():

Before:

```
bsz, seqlen, _ = x.shape
xq, xk, xv = self.wq(x), self.wk(x), self.wv(x)

xq = xq.view(bsz, seqlen, self.n_local_heads, self.head_dim)
xk = xk.view(bsz, seqlen, self.n_local_kv_heads, self.head_dim)
xv = xv.view(bsz, seqlen, self.n_local_kv_heads, self.head_dim)

xq, xk = apply_rotary_emb(xq, xk, freqs_cis=freqs_cis)

if self.kv_caching:
    self.cache_k = self.cache_k.to(xq)
    self.cache_v = self.cache_v.to(xq)

    self.cache_k[:,bsz, start_pos : start_pos + seqlen] = xk
    self.cache_v[:,bsz, start_pos : start_pos + seqlen] = xv

    keys = self.cache_k[:,bsz, : start_pos + seqlen]
    values = self.cache_v[:,bsz, : start_pos + seqlen]
else:
    pass
```

After:

```
bsz, seqlen, _ = x.shape
xq, xk, xv = self.wq(x), self.wk(x), self.wv(x)

xq = xq.view(bsz, seqlen, self.n_local_heads, self.head_dim)
xk = xk.view(bsz, seqlen, self.n_local_kv_heads, self.head_dim)
xv = xv.view(bsz, seqlen, self.n_local_kv_heads, self.head_dim)

xq, xk = apply_rotary_emb(xq, xk, freqs_cis=freqs_cis)

if self.kv_caching:
    self.cache_k = self.cache_k.to(xq)
    self.cache_v = self.cache_v.to(xq)

    self.cache_k[:,bsz, start_pos : start_pos + seqlen] = xk
    self.cache_v[:,bsz, start_pos : start_pos + seqlen] = xv

    keys = self.cache_k[:,bsz, : start_pos + seqlen]
    values = self.cache_v[:,bsz, : start_pos + seqlen]
else:
    keys = xk
    values = xv
```

2. In generation.py modify Generation.generate():

Before:

```
for cur_pos in range(min_prompt_len, total_len):
    with torch.no_grad():
        if kv_caching:
            logits = self(tokens[:, prev_pos:cur_pos], prev_pos)
        else:
            pass
        if temperature > 0:
            probs = torch.softmax(logits[:, -1] / temperature, dim=-1)
            next_token = sample_top_p(probs, top_p)
        else:
            next_token = torch.argmax(logits[:, -1], dim=-1)
```

After:

```
for cur_pos in range(min_prompt_len, total_len):
    with torch.no_grad():
        if kv_caching:
            logits = self(tokens[:, prev_pos:cur_pos], prev_pos)
        else:
            logits = self(tokens[:, : cur_pos], 0)
        if temperature > 0:
            probs = torch.softmax(logits[:, -1] / temperature, dim=-1)
            next_token = sample_top_p(probs, top_p)
        else:
            next_token = torch.argmax(logits[:, -1], dim=-1)
```

Steps taken to implement Gradient Accumulation:

1. In finetuning.py implemented the following code:

```
SET_GRADIENT_ACCUMULATION = bool(args.ga) # Set to True to enable gradient accumulation
if SET_GRADIENT_ACCUMULATION:
    ACCUMULATION_STEPS = 8 # Set to 8 for gradient accumulation
else:
    ACCUMULATION_STEPS = 1

if ((step + 1) % ACCUMULATION_STEPS == 0) or (step + 1 == len(dataloader)):
    if SET_MIXED_PRECISION:
        scaler.step(optimizer)
        scaler.update()
    else:
        optimizer.step()

    optimizer.zero_grad()
```

## Steps taken to implement Mixed Precision:

1. In `finetuning.py` implemented the following code:

```
SET_MIXED_PRECISION = bool(args.mp)
scaler = GradScaler('cuda') if SET_MIXED_PRECISION else None

if SET_MIXED_PRECISION:
    with autocast('cuda', dtype=torch.float16):
        # Forward pass - using the custom Llama model interface
        logits = model(tokens=input_ids, start_pos=0)

        # Calculate loss manually since the model doesn't handle it
        # Shift logits and labels for next token prediction
        shift_logits = logits[:, :-1, :].contiguous()
        shift_labels = labels[:, 1:].contiguous()

        # Flatten the tokens
        loss_fct = torch.nn.CrossEntropyLoss(ignore_index=IGNORE_INDEX)
        shift_logits = shift_logits.view(-1, shift_logits.size(-1))
        shift_labels = shift_labels.view(-1)

        shift_labels = shift_labels.to(shift_logits.device)
        loss = loss_fct(shift_logits, shift_labels)

    # Backward pass and optimization
    loss = loss / ACCUMULATION_STEPS
    scaler.scale(loss).backward() # Scale the loss for mixed precision
```

## Steps taken to implement LoRA:

1. In `lora.py` created a custom **LoRALayer Base Class**:
  - a. Initializes LoRA-specific parameters: rank ( $r$ ), scaling factor (`lora_alpha`), dropout (`lora_dropout`), and a flag to indicate if weights should be merged (`merge_weights`).
2. In `lora.py` created a **LoRALinear Class**:
  - a. Inherits from both `nn.Linear` and `LoRALayer`.
  - b. Introduces two trainable matrices, `lora_A` and `lora_B`, which represent the low-rank decomposition.
  - c. Overrides the forward method to add the LoRA adjustment to the original linear transformation.
3. In `lora.py` created a custom **mark\_only\_lora\_as\_trainable** function:
  - a. Freezes all model parameters except for `lora_A` and `lora_B`, ensuring that only LoRA parameters are updated during fine-tuning.
4. In `lora.py` created a custom **apply\_lora\_to\_llama** function:
  - a. Applies LoRA to the LLaMA model by replacing the query (`wq`) and value (`wv`) projection layers in each attention block with `LoRALinear` layers.
  - b. Copies the original weights to the new LoRA layers to maintain the pre-trained knowledge.
  - c. Calls `mark_only_lora_as_trainable` to freeze non-LoRA parameters.

5. In finetuning.py implemented the following code:

```
SET_LORA = bool(args.lora)

if SET_LORA:
    # Apply LoRA to the model
    model, lora_params_info = apply_lora_to_llama(
        model,
        rank=LORA_R,
        alpha=LORA_ALPHA,
        dropout=LORA_DROPOUT
    )
else:
    lora_params_info = (0, 0, 0) # Dummy values if LoRA is not used
```

Steps taken to implement Gradient Checkpointing:

1. In finetuning.py implemented the following code:

```
SET_CHECKPOINT = bool(args.gc)

model_args.use_gradient_checkpoint = SET_CHECKPOINT # Enable gradient checkpointing if specified
```

2. In model.py implemented the following code:

```
for i, layer in enumerate(self.layers):
    if self.params.use_gradient_checkpoint and i % 3 != 0:
        h = checkpoint(layer, h, start_pos, freqs_cis, mask)
    else:
        h = layer(h, start_pos, freqs_cis, mask)
```

**Justification for the gradient checkpointing implementation:**

- **Selective Layer Checkpointing:** Instead of checkpointing all layers, I'm applying it to 2 out of every 3 layers. This provides a good balance between memory savings and computational overhead.
- **Targeting Transformer Layers:** The transformer layers (self-attention + MLP) consume the most memory during training, so they're the best targets for checkpointing.
- **Preserving Original Structure:** The implementation preserves the model's original structure and API, making it compatible with the rest of your training code.
- **Compatible with LoRA:** This implementation works alongside LoRA fine-tuning, as checkpointing happens at the layer level while LoRA modifies individual weight matrices.
- **Memory reduction:** Approximately 30-40% less memory usage during training
- **Computation increase:** About 20-30% increase in training time due to recomputation
- **Better batch sizes:** Ability to use larger batch sizes or sequence lengths

By not checkpointing every layer, we maintain some computational efficiency while still getting significant memory savings.

### Performance Analysis:

		Grad. Accumulation	Grad. Checkpoint	Mixed Precision	LoRA
Memory	parameter	—	—	—	↓
	activation	↓	↓	↓	—
	gradient	↑	↑	↓	↓
	optimizer state	↑	—	↓	↓
Computation	FLOPs	—	↑	—	—
	runtime	↑	↑	↓	—

### Performance Benchmark with Gradient Accumulation:

GC	OFF				ON			
MP	OFF		ON		OFF		ON	
LoRA	OFF	ON	OFF	ON	OFF	ON	OFF	ON
<b>Peak Mem (MB)</b>	13993.84	7523.20	x	9630.22	13358.04	6864.00	x	7958.86
<b>Runtime (secs)</b>	38.02	22.64	x	27.00	40.99	29.13	x	39.57

### Performance Benchmark without Gradient Accumulation:

GC	OFF				ON			
MP	OFF		ON		OFF		ON	
LoRA	OFF	ON	OFF	ON	OFF	ON	OFF	ON
<b>Peak Mem (MB)</b>	11624.06	7521.13	11624.91	9620.55	11623.09	6857.48	11623.79	7954.38
<b>Runtime (secs)</b>	35.79	22.71	50.50	27.07	41.23	28.48	57.52	34.54

### Training Loss Over Time:

Note: When CUDA is out of memory while training a model it will output “OOM”.

#### - Vanilla Model:

Settings: Batch Size = 1, Epoch = 1, Learning rate =  $1 \times 10^{-5}$

Training Loss

Step 0: 3.0115  
 Step 10: 2.7962  
 Step 20: 2.6316  
 Step 30: 2.5735  
 Step 40: 2.6537  
 Step 50: 2.7595  
 Step 60: 2.5368  
 Step 70: 2.4725  
 Step 80: 2.3645  
 Step 90: 2.2995  
 Step 100: 2.2325  
 Step 110: 2.1597  
 Step 120: 2.0793  
 Step 130: 2.0061

Step 140: 2.0145  
Step 150: 1.9785  
Step 160: 1.9754  
Step 170: 1.9536  
Step 180: 1.9189  
Step 190: 1.8848  
Avg Loss: 1.8848  
Training time: 35.79 seconds  
Peak memory usage: 11624.06 MB  
Total parameters: 1498482688

- **Implemented Gradient Accumulation:**

Settings: Batch Size = 1, Epoch = 1, Learning rate =  $1 \times 10^{-5}$ , Gradient Accumulation\_Steps = 8

Training Loss  
Step 0: 0.3207  
Step 10: 0.3524  
Step 20: 0.3941  
Step 30: 0.3887  
Step 40: 0.3994  
Step 50: 0.3849  
Step 60: 0.3786  
Step 70: 0.3692  
Step 80: 0.3625  
Step 90: 0.3648  
Step 100: 0.3586  
Step 110: 0.3576  
Step 120: 0.3558  
Step 130: 0.3599  
Step 140: 0.3646  
Step 150: 0.3613  
Step 160: 0.3717  
Step 170: 0.3729  
Step 180: 0.3670  
Step 190: 0.3632  
Avg Loss: 0.3632  
Training time: 38.02 seconds  
Peak memory usage: 13993.84 MB  
Total parameters: 1498482688

- **Implemented Gradient Accumulation and Mixed Precision**

Settings: Batch Size = 1, Epoch = 1, Learning rate =  $1 \times 10^{-5}$ , Gradient Accumulation Steps = 8

OOM occurred during training.  
Gradient Accumulation: True  
Gradient Checkpointing: False  
Mixed Precision: True  
LoRA: False  
Batch Size: 1  
Learning Rate: 1e-05  
Epochs: 1  
Gradient Accumulation Steps: 8

- **Implemented Gradient Accumulation, Mixed Precision, LoRA**

Settings: Batch Size = 1, Epoch = 1, Learning rate =  $1 \times 10^{-5}$ , Gradient Accumulation Steps = 8, LoRA setting:  $r = 16$ ,  $\alpha = 32$  and  $dropout = 0.05$

Training Loss  
Step 0: 0.5352  
Step 10: 0.4834  
Step 20: 0.4283  
Step 30: 0.4011  
Step 40: 0.4066

Step 50: 0.3967  
Step 60: 0.4096  
Step 70: 0.3964  
Step 80: 0.3968  
Step 90: 0.3908  
Step 100: 0.3867  
Step 110: 0.3821  
Step 120: 0.3786  
Step 130: 0.3828  
Step 140: 0.3937  
Step 150: 0.3896  
Step 160: 0.3873  
Step 170: 0.3882  
Step 180: 0.3867  
Step 190: 0.3837  
Avg Loss: 0.3837  
Training time: 27.00 seconds  
Peak memory usage: 9630.22 MB  
Trainable parameters: 1703936, Total parameters: 1500186624  
Percentage of trainable parameters: 0.11358160196474329%

- **Implemented Gradient Accumulation, Mixed Precision, LoRA, Gradient Checkpointing**

Settings: Batch Size = 1, Epoch = 1, Learning rate =  $1 \times 10^{-5}$ , Gradient Accumulation Steps = 8, LoRA setting:  $r = 16$ ,  $\alpha = 32$  and  $dropout = 0.05$

Training Loss:  
Step 0: 0.4148  
Step 10: 0.3691  
Step 20: 0.3391  
Step 30: 0.3443  
Step 40: 0.3426  
Step 50: 0.3665  
Step 60: 0.3745  
Step 70: 0.3826  
Step 80: 0.3898  
Step 90: 0.3836  
Step 100: 0.3826  
Step 110: 0.3829  
Step 120: 0.3845  
Step 130: 0.3808  
Step 140: 0.3775  
Step 150: 0.3802  
Step 160: 0.3759  
Step 170: 0.3749  
Step 180: 0.3768  
Step 190: 0.3837  
Avg Loss: 0.3837  
Training time: 39.57 seconds  
Peak memory usage: 7958.86 MB  
Trainable parameters: 1703936, Total parameters: 1500186624  
Percentage of trainable parameters: 0.11358160196474329%

## Test prompts provided to all models:

```
prompts = [  
    # For these prompts, the expected answer is the natural continuation of the  
    prompt  
    "I believe the meaning of life is",  
    "Simply put, the theory of relativity states that ",  
    """"A brief message congratulating the team on the launch:  
  
    Hi everyone,  
  
    I just """,  
    # Few shot prompt (providing a few examples before asking model to complete  
    more);  
    """"Translate English to French:  
  
    sea otter => loutre de mer  
    peppermint => menthe poivrée  
    plush girafe => girafe peluche  
    cheese =>""",  
]
```

## Output Before Finetuning (Vanilla Llama Model):

```
I believe the meaning of life is  
> to live it, to taste experience to the utmost, to reach out eagerly and without  
fear for newer and richer experience. This is the meaning of life. To live it is to  
know a little of the meaning of life. To reach out eagerly is to know a little of  
the meaning of life. To taste experience is  
  
=====
```

Simply put, the theory of relativity states that  
> 1) all observers moving with respect to each other will measure the same length  
of a meter stick; 2) all observers moving with respect to each other will measure  
the same speed of light; 3) all observers moving with respect to each other will  
measure the same time interval between two events. This means that the

```
=====
```

A brief message congratulating the team on the launch:

```
Hi everyone,  
  
I just  
> received the notification that the team has successfully launched the app.  
I'm very proud of the team and the work they have put in to make this a  
reality.  
  
I hope that the app will help you all to learn more about the world of  
Cryptocurrency and the blockchain technology.  
  
I wish you
```



=====

Translate English to French:

```
    sea otter => loutre de mer
    peppermint => menthe poivrée
    plush girafe => girafe peluche
    cheese =>
> fromage
    car => voiture
    toot => trombone
    pipsqueak => petit chat
    zephyr => vent léger
    sardine => sardine
    tramp => marchand
    aardvark => ariette
    pretzel => p
```

=====

Model	Output after finetuning
Vanilla Model	<p>I believe the meaning of life is</p> <p>&gt; to find your passion and do what you love. I have a passion for photography, writing, and traveling. I love to learn and experience new things. I am always open to new adventures and have a passion for meeting new people. I have a love for the outdoors and enjoy spending time outdoors. I am a hard</p> <p>=====</p> <p>Simply put, the theory of relativity states that</p> <p>&gt; 1) The speed of light is the same for all observers, regardless of their motion relative to the source of light, and 2) The laws of physics are the same in all reference frames. The theory of relativity states that the laws of physics are the same in all reference frames. This means that the same</p> <p>=====</p> <p>A brief message congratulating the team on the launch:</p> <p>Hi everyone,</p> <p>I just</p> <p>&gt; launched the website for the first time!</p> <p>Please check it out at: <a href="https://www.kcduke.com/">https://www.kcduke.com/</a></p> <p>Best wishes,</p> <p>Kevin</p> <p>=====</p> <p>Translate English to French:</p> <pre>    sea otter =&gt; loutre de mer     peppermint =&gt; menthe poivrée     plush girafe =&gt; girafe peluche     cheese =&gt; &gt; fromage     fennel =&gt; fenouil     tangerine =&gt; orange tangerine     apple =&gt; pomme     tuna =&gt; saumon</pre>

	<p>pineapples =&gt; pommes de pin  fish =&gt; poisson  toffee =&gt; confiture au theï  goat =&gt;</p> <p>=====</p>
<p>Implemented:</p> <ul style="list-style-type: none"> <li>- Gradient Accumulation</li> <li>- Mixed Precision</li> <li>- LoRA</li> <li>- Gradient Checkpointing</li> </ul>	<p>I believe the meaning of life is  &gt; to live it, to taste experience to the utmost, to reach out eagerly and without fear for newer and richer experience. This is the meaning of life. To live it is to know a little of I and a little of Thou, and to love a little of both.  I am a Canadian born and raised in the</p> <p>=====</p> <p>Simply put, the theory of relativity states that  &gt; 1) all observers moving with respect to each other will measure the same length of a meter stick; 2) all observers moving with respect to each other will measure the same speed of light; 3) all observers moving with respect to each other will measure the same time interval between two events. This means that the</p> <p>=====</p> <p>A brief message congratulating the team on the launch:</p> <p>Hi everyone,</p> <p>I just  &gt; received the notification that the team has successfully launched the app.  I'm very proud of the team and the work they have put in to make this a reality.</p> <p>I hope that the app will help you all to learn more about the world of investment, and to become more confident in your own financial future</p> <p>=====</p> <p>Translate English to French:</p> <p>sea otter =&gt; loutre de mer  peppermint =&gt; menthe poivrée  plush girafe =&gt; girafe peluche  cheese =&gt;  &gt; fromage  car =&gt; voiture  toot =&gt; trombone  pipsqueak =&gt; petit chat  zephyr =&gt; vent léger  sardine =&gt; sardine  yukon =&gt; Yukon  prairie =&gt; prairie  gnat =&gt; grenouille</p> <p>=====</p>